

Project - Online News Popularity

[Code](#)[Hide](#)

```
rm(list = ls())
```

In this study, we need the following packages

[Hide](#)

```
library(tidyverse)
library(glmnet)
library(leaps)
library(randomForest)
theme_set(theme_bw())
```

1. Description of the practical problem

Mashable is a global, multi-platform media and entertainment company. The following dataset summarizes a heterogeneous set of features about articles published by Mashable in a period of two years. The database can be found at the following address: <https://archive.ics.uci.edu/ml/datasets/online+news+popularity> (<https://archive.ics.uci.edu/ml/datasets/online+news+popularity>)

We read the data.

[Hide](#)

```
mydata = read.csv(file = "OnlineNewsPopularity.csv", header=TRUE, sep=";", na.strings = "NA")
head(mydata, n = 1)
```

url <fctr>	timedelta <dbl>	n_tokens_title <dbl>
1 http://mashable.com/2013/01/07/amazon-instant-video-browser/	731	12

1 row | 1-4 of 61 columns

[Hide](#)

```
dim(mydata)
```

```
[1] 39644 61
```

Our data set is quite large. We have 39644 observations and 61 variables. The goal is to predict the number of shares on social networks (popularity).

2. Description of the mathematical problem

The problem is to explain the output that will be named "nb_shares" by the other variables of the dataset (we will exclude some of them). We denote by Y the dependent variable "nb_shares" and by X_n the n independent variables. We want to explain the connection between Y and X_1, \dots, X_n .

Y being continuous, we are facing a regression problem. We want to find a machine that will minimize the error between one prediction of our model y' and one observation Y . In order to measure this error, we will use the **quadratic risk function** for the machine m that is $R(m) = E[(Y - m(X))^2]$ to try to find a machine m such as $m_n(x) \approx m^*(x)$.

As we will see later, we will estimate these risks by splitting our dataset into a **train** dataset and **test** dataset.

3. Preparation and cleaning of the data

Cleaning of the names of the columns

- Transform lower cases

[Hide](#)

```
colnames(mydata) = tolower(colnames(mydata))
```

- **Change the columns names to have more explicit names**

Hide

```
colnames(mydata)[colnames(mydata)=="n_tokens_title"] = "nb_words_title"
colnames(mydata)[colnames(mydata)=="n_tokens_content"] = "nb_words_content"
colnames(mydata)[colnames(mydata)=="n_unique_tokens"] = "rate_unique_words"
colnames(mydata)[colnames(mydata)=="n_non_stop_words"] = "rate_non_stop_words"
colnames(mydata)[colnames(mydata)=="n_non_stop_unique_tokens"] = "rate_non_stop_unique_words"
colnames(mydata)[colnames(mydata)=="num_hrefs"] = "nb_links"
# Number of links to other articles published by Mashable
colnames(mydata)[colnames(mydata)=="num_self_hrefs"] = "nb_links_mashable"
colnames(mydata)[colnames(mydata)=="num_imgs"] = "nb_images"
colnames(mydata)[colnames(mydata)=="num_videos"] = "nb_videos"
colnames(mydata)[colnames(mydata)=="average_token_length"] = "average_word_length"
colnames(mydata)[colnames(mydata)=="num_keywords"] = "nb_keywords"
colnames(mydata)[colnames(mydata)=="shares"] = "nb_shares" # Target variable
```

- **Cleaning the data**

We have ratio superior to 1. It is not possible. We find the row where the ratio of rate_non_stop_words is maximum.

Hide

```
index_row = which.max(mydata$rate_non_stop_words)
mydata[index_row,]
```

url	timedelta	nb_words_tit
<fctr>	<dbl>	<dbl>
31038 http://mashable.com/2014/08/18/ukraine-civilian-convoy-attacked/	142	
1 row 1-4 of 61 columns		

A lot of aberrant values. We drop this row.

Hide

```
mydata = mydata[-index_row,]
```

We also detect an article with 0 words. We check it.

Hide

```
index_row = which.min(mydata$nb_words_content)
mydata[index_row,]
```

url	timedelta	nb_words_title
<fctr>	<dbl>	<dbl>
894 http://mashable.com/2013/01/23/actual-facebook-graph-searches/	715	10
1 row 1-4 of 61 columns		

This is a mistake there are words in this article. We exclude the rows without words.

Hide

```
dim(mydata) # Before exclusion
```

```
[1] 39643    61
```

Hide

```
mydata = mydata %>% filter(nb_words_content!=0)
dim(mydata) # After exclusion
```

```
[1] 38462    61
```

We drop the column `rate_non_stop_words` with only 1 in values.

Hide

```
mydata = mydata %>% select(-rate_non_stop_words)
```

- **NA values**

We check that we do not have NA values. Look at the number by column.

Hide

```
na_columns = sapply(mydata, function(x) sum(is.na(x)))
na_columns[na_columns>0]
```

```
named integer(0)
```

We do not have any missing values.

- **Creation of categorical variables**

0 - Create an ID by row

Hide

```
number_rows = dim(mydata)[1]
ID = data.frame(ID = c(1:number_rows))
mydata = cbind(ID,mydata)
```

Transforming the categorical variables to factor.

Hide

```
columns_weekdays = mydata %>% select(contains("weekday")) %>% colnames
columns_channel = mydata %>% select(contains("data_channel_is")) %>% colnames
other_columns = c("is_weekend")
# All the columns names in one vector
columns_to_factor = c(columns_weekdays,columns_channel,other_columns)
# We change them to factor.
mydata[columns_to_factor] = lapply(mydata[columns_to_factor], factor)
```

- **Group the days and the channel in two variables.**

We have 7 categorical variables: `weekday_is_monday` and so on. We create a new categorical variable "weekday" with the name of the day.

Hide

```
df_weekday = mydata %>% select(ID, columns_weekdays) %>% gather(key=weekday, value=value_to_filter, -ID)
df_weekday = df_weekday %>% mutate(weekday=recode(weekday,

"weekday_is_monday"="monday",
"weekday_is_tuesday"="tuesday",

"weekday_is_wednesday"="wednesday",
"weekday_is_thursday"="thursday",

"weekday_is_friday"="friday",
"weekday_is_saturday"="saturday",

"weekday_is_sunday"="sunday"
))

df_weekday = df_weekday %>% filter(value_to_filter==1)
df_weekday$weekday = as.factor(df_weekday$weekday)
df_weekday$weekday = ordered(df_weekday$weekday, levels=c("monday", "tuesday", "wednesday", "thursday", "friday", "saturday", "sunday"))
df_weekday = df_weekday %>% select(ID, weekday)
mydata = left_join(mydata, df_weekday, by="ID")
```

We have 6 categorical variables: `data_channel_is_lifestyle` and so on. We create a new categorical variable "data_channel" with the name of the channel.

Hide

```
df_channel = mydata %>% select(ID, columns_channel) %>% gather(key=channel, value=value_to_filter, -ID)
df_channel = df_channel %>% mutate(channel=recode(channel,
"lifestyle",
"data_channel_is_lifestyle"=
"entertainment",
"data_channel_is_entertainmen
"med",
"data_channel_is_bus"="bus",
"data_channel_is_socmed"="soc
,
"data_channel_is_tech"="tech"
d"
"data_channel_is_world"="worl
))

df_channel = df_channel %>% filter(value_to_filter==1)
df_channel$channel = as.factor(df_channel$channel)
levels(df_channel$channel) = c(levels(df_channel$channel),"other")
df_channel = df_channel %>% select(ID, channel)
mydata = left_join(mydata, df_channel,by="ID")
mydata$channel[is.na(mydata$channel)] = "other"
```

- **Removal of non predictive variables**

We move the target variable to first position after the non predictive variables.

Hide

```
mydata = mydata %>% select(url, ID, timedelta, nb_shares, everything())
```

To create the models, we create a data.frame "mydata2" without the doublons and the non predictive variables.

Hide

```
mydata2 = mydata %>% select(-c(columns_channel, columns_weekdays,url, ID, timedelta))
```

4. Description of the dataset (with dplyr)

Hide

```
summary(mydata2)
```


Median :0.04000	Median :0.05000	Median :0.4566	Median : 0.12252	Median :0.03960
Mean :0.21430	Mean :0.23870	Mean :0.4570	Mean : 0.12298	Mean :0.04084
3rd Qu.:0.34052	3rd Qu.:0.41461	3rd Qu.:0.5103	3rd Qu.: 0.17992	3rd Qu.:0.05072
Max. :0.92653	Max. :0.92719	Max. :1.0000	Max. : 0.72784	Max. :0.15549

global_rate_negative_words	rate_positive_words	rate_negative_words	avg_positive_polarity	min_positive_polarity
Min. :0.00000	Min. :0.0000	Min. :0.0000	Min. :0.0000	Min. :0.00
1st Qu.:0.01018	1st Qu.:0.6129	1st Qu.:0.2000	1st Qu.:0.3119	1st Qu.:0.05
Median :0.01567	Median :0.7143	Median :0.2857	Median :0.3619	Median :0.10
Mean :0.01712	Mean :0.7031	Mean :0.2968	Mean :0.3647	Mean :0.09
3rd Qu.:0.02199	3rd Qu.:0.8000	3rd Qu.:0.3871	3rd Qu.:0.4133	3rd Qu.:0.10
Max. :0.18493	Max. :1.0000	Max. :1.0000	Max. :1.0000	Max. :1.00

max_positive_polarity	avg_negative_polarity	min_negative_polarity	max_negative_polarity	title_subjectivity
Min. :0.00	Min. : -1.0000	Min. : -1.0000	Min. : -1.0000	Min. :0.000
1st Qu.:0.60	1st Qu.: -0.3315	1st Qu.: -0.7143	1st Qu.: -0.1250	1st Qu.:0.000
Median :0.80	Median : -0.2577	Median : -0.5000	Median : -0.1000	Median :0.125
Mean :0.78	Mean : -0.2675	Mean : -0.5380	Mean : -0.1108	Mean :0.280
3rd Qu.:1.00	3rd Qu.: -0.1934	3rd Qu.: -0.3125	3rd Qu.: -0.0500	3rd Qu.:0.500
Max. :1.00	Max. : 0.0000	Max. : 0.0000	Max. : 0.0000	Max. :1.000

title_sentiment_polarity	abs_title_subjectivity	abs_title_sentiment_polarity	weekday
Min. : -1.0000	Min. :0.0000	Min. :0.0000	monday :6471 bus
1st Qu.:0.0000	1st Qu.:0.1667	1st Qu.:0.0000	tuesday :7170 enterta
Median :0.0000	Median :0.5000	Median :0.0000	wednesday:7205 lifesty
Mean :0.0710	Mean :0.3424	Mean :0.1549	thursday :7052 socmed
3rd Qu.:0.1364	3rd Qu.:0.5000	3rd Qu.:0.2500	friday :5538 tech
Max. :1.0000	Max. :0.5000	Max. :1.0000	saturday :2369 world
			sunday :2657 other

[Hide](#)

```
colnames(mydata2)
```

[1] "nb_shares"	"nb_words_title"	"nb_words_content"	"rat
e_unique_words"			
[5] "rate_non_stop_unique_words"	"nb_links"	"nb_links_mashable"	"nb_
images"			
[9] "nb_videos"	"average_word_length"	"nb_keywords"	"kw_
min_min"			
[13] "kw_max_min"	"kw_avg_min"	"kw_min_max"	"kw_
max_max"			
[17] "kw_avg_max"	"kw_min_avg"	"kw_max_avg"	"kw_
avg_avg"			
[21] "self_reference_min_shares"	"self_reference_max_shares"	"self_reference_avg_shares"	"is_
weekend"			
[25] "lda_00"	"lda_01"	"lda_02"	"lda
_03"			
[29] "lda_04"	"global_subjectivity"	"global_sentiment_polarity"	"glo
bal_rate_positive_words"			
[33] "global_rate_negative_words"	"rate_positive_words"	"rate_negative_words"	"avg
_positive_polarity"			
[37] "min_positive_polarity"	"max_positive_polarity"	"avg_negative_polarity"	"min
_negative_polarity"			
[41] "max_negative_polarity"	"title_subjectivity"	"title_sentiment_polarity"	"abs
_title_subjectivity"			
[45] "abs_title_sentiment_polarity"	"weekday"	"channel"	

Columns numeric

Hide

```
columns_numeric <- unlist(lapply(mydata2, is.numeric))
```

- We follow our intuition and look at some correlations that could be interesting

Hide

```
cor(mydata2[,c("nb_shares", "nb_words_title", "nb_words_content", "nb_images", "nb_videos")])
```

	nb_shares	nb_words_title	nb_words_content	nb_images	nb_videos
nb_shares	1.000000000	0.006212729	0.006701789	0.041279165	0.02471476
nb_words_title	0.006212729	1.000000000	0.028162440	-0.006525119	0.05246549
nb_words_content	0.006701789	0.028162440	1.000000000	0.352948886	0.10205617
nb_images	0.041279165	-0.006525119	0.352948886	1.000000000	-0.06657575
nb_videos	0.024714759	0.052465492	0.102056168	-0.066575748	1.00000000

The correlations are lower than we expected ! Having a lot of content in an article does not make it always popular.

- Overview - Main features of an article

Hide

```
mydata2%>%select(Length=title=nb_words_title,Nb_Links=nb_links,Subjectivity=global_subjectivity,Positive=max_positive_polarity,Content=nb_words_content)%>%summarise_all(funs(mean))
```

Length title <dbl>	Nb_Links <dbl>	Subjectivity <dbl>	Positive <dbl>	Content <dbl>
10.38246	11.21788	0.4569957	0.779983	563.2692
1 row				

- Does adding multimedia content have an impact on number of shares?

Hide

```
mydata2%>%select(nb_images,nb_videos,nb_shares)%>%mutate(multimedia_content=nb_images+nb_videos)%>%group_by(multimedia_content)%>%summarise(mean_shares=mean(nb_shares))%>%arrange(desc(mean_shares))%>%filter(mean_shares<10000)%>%head()
```

multimedia_content <dbl>	mean_shares <dbl>
-----------------------------	----------------------

multimedia_content <dbl>	mean_shares <dbl>
55	9588.375
71	9437.000
90	9400.000
46	8450.000
60	8243.667
58	7500.000
6 rows	

Articles with the higher shares mean were article with lots of multimedia content (excluding articles considered as outliers with more than 10 000 shares).

- Does an article with a subjective title drive more shares?

Hide

```
subjectivity=mydata2%>%select("nb_shares", "title_subjectivity")%>%filter(nb_shares<20000)
subjectivity$title_subjectivity=as.factor(as.numeric(subjectivity$title_subjectivity>0.75))
subjectivity%>%select(title_subjectivity,nb_shares)%>%group_by(title_subjectivity)%>%summarise(Shares=
mean(nb_shares))%>%arrange(desc(Shares))
```

title_subjectivity <fctr>	Shares <dbl>
1	2838.307
0	2379.921
2 rows	

An article with subjectivity in its title tends to drive slightly more shares in average.

5. Vizualization of the dataset (with ggplot)


- How the number of shares is distributed?

Hide

```
ggplot(mydata2) + aes(x=nb_shares) + geom_density(fill="darkblue") + xlim(0,20000)+xlab("Number of sha
res")+ylab("Density")+ labs(title = "Shares Density", subtitle = "Excluding articles with more than 20
000 shares") +
  theme(plot.title = element_text(hjust = 0.5,face="bold"),plot.subtitle = element_text(hjust = 0.5))
```


Shares Density

Excluding articles with more than 20000 shares



We observe that most of the articles have a number of shares less than 5000.

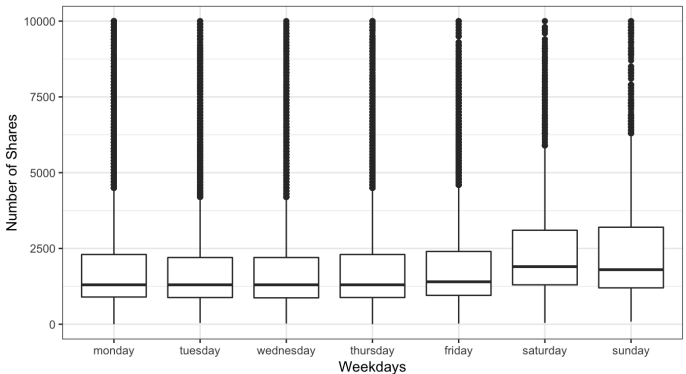
- Which weekday drive more shares?

Hide

```
ggplot(mydata2) + aes(x=weekday, y=nb_shares) + geom_boxplot() + ylim(0, 10000)+xlab("Weekdays")+ylab("Number of Shares")+ labs(title = "Which day drive more shares?", subtitle = "Excluding articles with more than 10000 shares") + theme(plot.title = element_text(hjust = 0.5,face="bold"),plot.subtitle = element_text(hjust = 0.5))
```

Which day drive more shares?

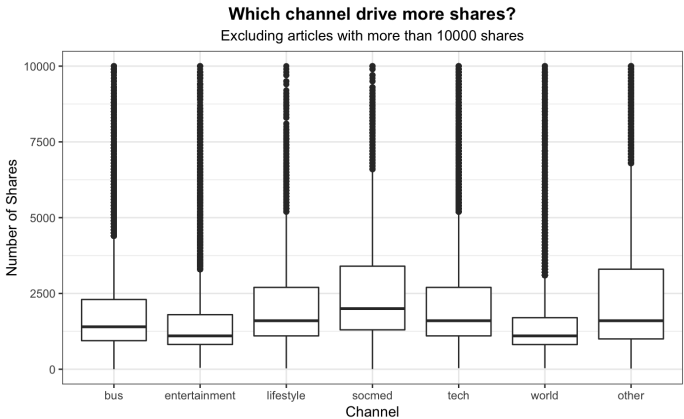
Excluding articles with more than 10000 shares



We can observe that Saturday and Sunday drive more shares in average than the other weekdays.

- Which channel drive more shares?

```
ggplot(mydata2) + aes(x=channel, y=nb_shares) + geom_boxplot() + ylim(0, 10000)+xlab("Channel")+ylab("Number of Shares")+ labs(title = "Which channel drive more shares?", subtitle = "Excluding articles with more than 10000 shares") +
  theme(plot.title = element_text(hjust = 0.5,face="bold"),plot.subtitle = element_text(hjust = 0.5))
```



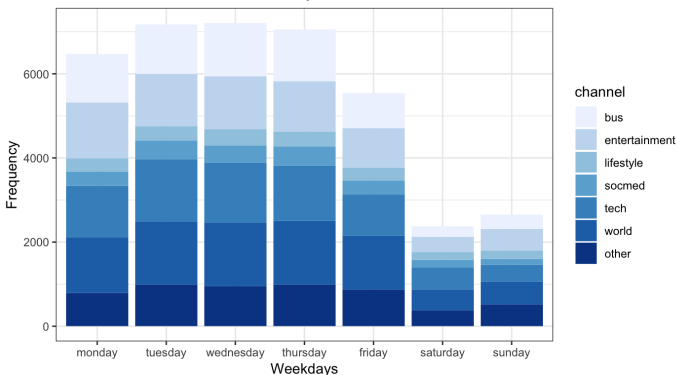
The above boxplots show that lifestyle, socmed, tech and other are the four channels that drive in average more shares as we have seen in the description of data.

- **Comparison between weekday & Channel**

```
ggplot(mydata2) + aes(x=weekday, fill=channel)+ geom_bar(position = "stack")+scale_fill_brewer(palette = "Blues")+xlab("Weekdays")+ylab("Frequency")+ labs(title = "Overview by day - Frequency of article posted", subtitle = "Detailed by channel") +
  theme(plot.title = element_text(hjust = 0.5,face="bold"),plot.subtitle = element_text(hjust = 0.5))
```

Overview by day - Frequency of article posted

Detailed by channel



We mentioned before that during weekend, articles that are posted drive more shares. Even though, we notice that there is less articles posted compared to other weekdays.

- **Is there a link between shares and word content?**

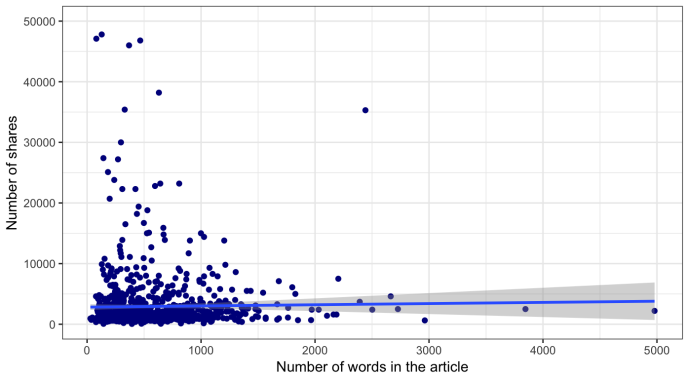
To do this plot we select a sub-sample.

Hide

```
mydata2_temp = sample_n(tbl = mydata2, size = 1000)
ggplot(mydata2_temp) + aes(x=nb_words_content,y=nb_shares)+ geom_point(color="darkblue") +
  geom_smooth(method="lm") +
  ylim(0,50000)+ylab("Number of shares")+xlab("Number of words in the article") +
  labs(title = "Is there a link between shares and word content?", subtitle = "Excluding articles with
more than 50000 shares") +
  theme(plot.title = element_text(hjust = 0.5,face="bold"),plot.subtitle = element_text(hjust = 0.5))
```

Is there a link between shares and word content?

Excluding articles with more than 50000 shares



The above graph suggests that the linear relation is weak. However we can see that most of the very shared articles do not have a lot of words.

6. Machine Learning Methods

We propose to use the following methods to predict the number of shares:

1 - Linear model

2 - Penalized Regression (Lasso and Ridge)

3 - Random Forest

We create a copy of that dataset that will be used for models

Hide

```
data <- mydata2
```

Additional preliminary cleaning

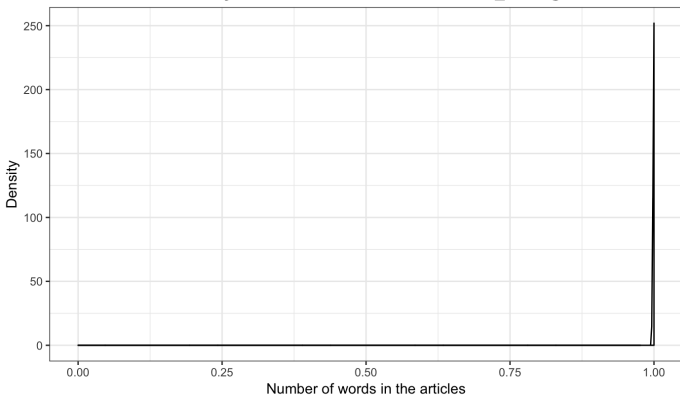
There are linear dependencies in the data that will prevent us from running some models :

rate_positive_words + rate_negative_words = 1

Hide

```
data %>% mutate(word_feeling = rate_positive_words + rate_negative_words) %>% ggplot() + aes(x = word_feeling) + geom_density() + ylab("Density") + xlab("Number of words in the articles") + labs(title = "Density of the sum of the variable word_feeling") + theme(plot.title = element_text(hjust = 0.5, face = "bold"))
```

Density of the sum of the variable word_feeling

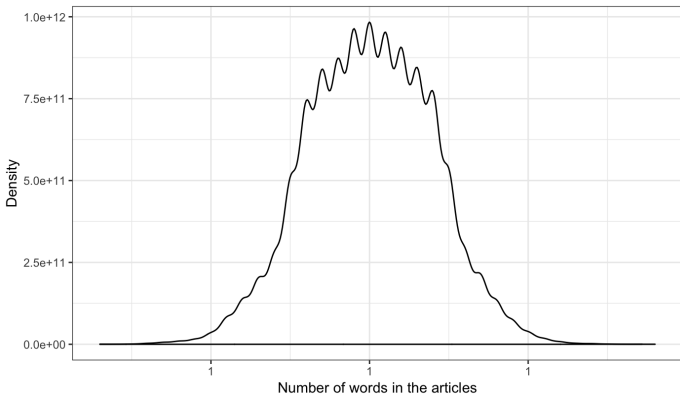


lda_00 + lda_01 + lda_02 + lda_03 + lda_04 = 1

Hide

```
data %>% mutate(sum_lda = lda_00 + lda_01 + lda_02 + lda_03 + lda_04) %>% ggplot() + aes(x = sum_lda)
+ geom_density() + ylab("Density")+xlab("Number of words in the articles") +
  labs(title = "Density of the sum of the variable word_feeling") + theme(plot.title = element_text(h
just = 0.5,face="bold"))
```

Density of the sum of the variable word_feeling



Hide

```
summary(data$lda_00 + data$lda_01 + data$lda_02 + data$lda_03 + data$lda_04)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
1	1	1	1	1	1

In abscciss, we have values of 1 due to error approximation in the sum. If we had truncated the values, it would have been only a peak at 1.

We exclude the two problematic variables.

Hide

```
data <- data %>% select(-c(lda_04,rate_negative_words))
```

Linear Model

We will try to predict the number of shares based on all other regressors We start naively by trying to fit a Linear Regression to the whole dataset

Hide

```
# model.lm <- lm(nb_shares~.,data=data)
# Error: cannot allocate vector of size 11.0 Gb
```

R cannot compute it, due to the excess of parameters We will therefore : - work only on a sample of the dataset - consider numerical values only

Hide

```
Name_numeric_columns <- rownames(data.frame(columns_numeric[columns_numeric]))
Name_numeric_columns <- Name_numeric_columns[Name_numeric_columns != "lda_04" & Name_numeric_columns != "rate_negative_words"]

set.seed(123456)
data <- data %>% sample_n(10000) %>% select(Name_numeric_columns)
```

We can now separate in training and testing datasets

Hide

```
data.train <- data %>% slice(1:8000)
data.test <- data %>% slice(8001:10000)
```

We can now run a naive prediction model (that predicts the sample mean of the target column in the train data) and a linear model

Hide


```
pred.naive <- rep(mean(data.train$nb_shares),nrow(data.test))
model.lm <- lm(nb_shares~.,data=data.train)
```

We do a subset selection. We select the **BIC** criteria rather than the **AIC** criteria because it penalises more aggressively the large models. Indeed, compared to **AIC**, **BIC** increases by $\log(n)d$ whereas **AIC** increases by $2d$. For $n = 10000$, we will have $\log(n)d \approx 9.21$ which is superior to 2 of the **AIC**.

Hide

```
reg.fit.forward <- regsubsets(nb_shares~.,data=data.train,method="forward",nvmax=150)
reg.fit.backward <- regsubsets(nb_shares~.,data=data.train,method="backward",nvmax=150)
sum.reg.fit.f <- summary(reg.fit.forward)
sum.reg.fit.b <- summary(reg.fit.backward)
nb.bic.f <- order(sum.reg.fit.f$bic)[1]
nb.bic.b <- order(sum.reg.fit.b$bic)[1]
get.formula<- function(a,number){
  var.sel <- a$which[number,][1]
  var.sell <- names(var.sel)[var.sel] %>% paste(collapse="+")
  form <- formula(paste("nb_shares~",var.sell,sep=""))
  str_form <- paste("nb_shares~",var.sell,sep="")
  str_form
}
plot(reg.fit.forward,scale="bic",main="Forward BIC")
```


Forward BIC



Hide

```
plot(reg.fit.backward,scale="bic",main="Backward BIC ")
```

Backward BIC



Hide

```
subset.models <- data.frame(Criteria=rep(c("Bic"),each=2),Method=c("Forward","Backward"),Formula = c(get.formula(sum.reg.fit.f,nb.bic.f),get.formula(sum.reg.fit.b,nb.bic.b)))
subset.models
```

Criteria	Method
<fctr>	<fctr>

Criteria <fctr>	Method <fctr>	
Bic	Forward	
Bic	Backward	
2 rows 1-2 of 3 columns		

We see that the variables selected with the forward or backward approaches are approximately the same.

For the forward approach, the BIC criteria selects the model:

$$Y = \beta_0 + nb.links * \beta_{nb.links} + kw.max.max * \beta_{kw.max.max} + kw.min.avg * \beta_{kw.min.avg} + kw.max.avg * \beta_{kw.max.avg} + kw.avg.avg * \beta_{kw.avg.avg} + self.reference.max.shares * \beta_{self.reference.max.shares} + lda.03 * \beta_{lda.03}$$

We evaluate those models based on the RMSE

Hide


```
RMSE.list <- c()
RMSE.list <- append(RMSE.list, mean((pred.naive-data.test$nb_shares)^2)^(1/2))
RMSE.list <- append(RMSE.list, mean((predict(model.lm, newdata=data.test)-data.test$nb_shares)^2)^(1/2))
for (formula in subset.models$Formula) {
  model.sub <- lm(formula,data=data.train)
  pred <- predict(model.sub, newdata = data.test)
  RMSE <- mean((pred-data.test$nb_shares)^2)^(1/2)
  RMSE.list <- append(RMSE.list,RMSE)
}
RMSE.comparison <- data.frame(Criteria="mean",Method="Naive",Formula="Y")
RMSE.comparison <- rbind(RMSE.comparison,data.frame(Criteria="LM",Method="LM",Formula="Y"),subset.mode
ls)
RMSE.comparison <- RMSE.comparison %>% select(Criteria,Method) %>% mutate(RMSE=RMSE.list)
RMSE.comparison
```

Criteria <fctr>	Method <fctr>	RMSE <dbl>
mean	Naive	5472.156
LM	LM	5378.744
Bic	Forward	5364.830
Bic	Backward	5372.931
4 rows		

Penalized regression


Hide

```
data.glmnet <- model.matrix(nb_shares~.,data=data.train)
model.ridge <- glmnet(data.glmnet,data.train$nb_shares,alpha=0)
model.lasso <- glmnet(data.glmnet,data.train$nb_shares,alpha=1)
plot(model.ridge)
```

```
plot(model.lasso)
```


Hide



We can observe the regularization path for Lasso and Ridge Regression. For Lasso, some $\beta = 0$ for small t .


```
ridgeCV <- cv.glmnet(data.glmnet, data.train$nb_shares, lambda=exp(seq(-20, 80, length=300)), alpha=0)
plot(ridgeCV)
```

Hide



Hide

```
lassoCV <- cv.glmnet(data.glmnet, data.train$nb_shares, lambda=exp(seq(-10, 20, length=300)), alpha=1)
plot(lassoCV)
```



We look at the 2 best models

Hide

```

model.ridge <- glmnet(data.glmnet,data.train$nb_shares,lambda=ridgeCV$lambda.min,alpha=0)
model.lasso <- glmnet(data.glmnet,data.train$nb_shares,lambda=lassoCV$lambda.min,alpha=1)
RMSE.ridge <- mean((predict(model.ridge, newx=model.matrix(nb_shares=.,data=data.test))-data.test$nb_shares)^2)^(1/2)
RMSE.lasso <- mean((predict(model.lasso, newx=model.matrix(nb_shares=.,data=data.test))-data.test$nb_shares)^2)^(1/2)
RMSE.penalized <- data.frame(Criteria=paste("Lambda = ",ridgeCV$lambda.min),Method="Ridge",RMSE=RMSE.ridge)
RMSE.penalized <- rbind(RMSE.penalized, data.frame(Criteria=paste("Lambda = ",lassoCV$lambda.min),Method="Lasso",RMSE=RMSE.lasso))
RMSE.comparison <- rbind(RMSE.comparison, RMSE.penalized)
RMSE.comparison

```

Criteria <fctr>	Method <fctr>	RMSE <dbl>
mean	Naive	5472.156
LM	LM	5378.744
Bic	Forward	5364.830
Bic	Backward	5372.931
Lambda = 4177539791803456000	Ridge	5472.156
Lambda = 46.8126678173277	Lasso	5357.719

6 rows

Random Forests

We reduce the number and size of tree to make it compute in a timely manner

Hide

```

model.rf <- randomForest(nb_shares=.,data=data.train,nodesize=30,ntree=300)
model.rf

```

```

Call:
randomForest(formula = nb_shares ~ ., data = data.train, nodesize = 30, ntree = 300)
Type of random forest: regression
Number of trees: 300
No. of variables tried at each split: 13

Mean of squared residuals: 59282165
% Var explained: -0.17

```

Hide

```

RMSE.rf <- mean((predict(model.rf, newdata=data.test)-data.test$nb_shares)^2)^(1/2)
RMSE.comparison <- rbind(RMSE.comparison, data.frame(Criteria="Node size = 30, ntree = 300",Method="RF",RMSE=RMSE.rf))
RMSE.comparison

```

Criteria <fctr>	Method <fctr>	RMSE <dbl>
mean	Naive	5472.156
LM	LM	5378.744
Bic	Forward	5364.830
Bic	Backward	5372.931
Lambda = 4177539791803456000	Ridge	5472.156
Lambda = 46.8126678173277	Lasso	5357.719

Criteria <fctr>	Method <fctr>	RMSE <dbl>
Node size = 30, ntree = 300	RF	5561.902
7 rows		

We increase parameters to get a more accurate prediction

Hide

```
model.rf2 <- randomForest(nb_shares=., data=data.train, nodesize=1, ntree=400)
RMSE.rf <- mean((predict(model.rf2, newdata=data.test)-data.test$nb_shares)^2)^(1/2)
RMSE.comparison <- rbind(RMSE.comparison, data.frame(Criteria="Node size = 1, ntree = 400", Method="RF"
, RMSE=RMSE.rf))
RMSE.comparison
```

Criteria <fctr>	Method <fctr>	RMSE <dbl>
mean	Naive	5472.156
LM	LM	5378.744
Bic	Forward	5364.830
Bic	Backward	5372.931
Lambda = 4177539791803456000	Ridge	5472.156
Lambda = 46.8126678173277	Lasso	5357.719
Node size = 30, ntree = 300	RF	5561.902
Node size = 1, ntree = 400	RF	5607.964
8 rows		


It doesn't improve the prediction, quite the contrary

7. Performances of each models

Based on our table of comparison of RMSE for each models, Lasso is the best performing model

Hide


```
pred <- predict(model.lasso, newx=model.matrix(nb_shares=., data=data.test))
obs <- data.test$nb_shares
res <- obs-pred
Prediction <- data.frame(pred, obs, res)
colnames(Prediction) <- c("pred", "obs", "res")
Prediction %>% ggplot() + aes(x=res) + geom_density()
```



Our model is quite optimistic (peak before 0 in residual's distribution), but still misses the few viral articles that get shared a lot. These are the ones that have a major impact on the poor RMSE score


Hide

```
Prediction <- Prediction %>% mutate(ratio = res/obs)  
Prediction %>% ggplot() + aes(x=ratio) + geom_density()
```



Hide


```
Prediction %>% ggplot() + aes(x=pred, y=obs, color=res) + geom_point()
```



We can see very few articles reach more than 12500 shares. They get very poor predictions as the model never predicts more than 11000 shares, and have a significant impact on the overall RMSE

Hide

```
Pred.trim <- Prediction[Prediction$obs < 12500,]
Pred.trim %>% ggplot() + aes(x=pred, y=obs, color=res) + geom_point()
```



Hide

```
mean(Pred.trim$res^2)^(1/2)
```

If we had only observations with less than 12500 shares in the test data, the RMSE could have had been halved. If on top of that we would not have viral articles in the train data, the RMSE would have been improved. However, because our goal was to predict the number of shares for all articles (including the ones with more than 12500 shares), we did not exclude popular articles from the database.

8. Conclusion

The popularity of an article depends on many characteristics: day of publication, length, channel, content... Using these characteristics as predictors, we have tried to predict the number of shares on social networks (popularity). Throughout analysis and modelling, we realise that there isn't a stable and standard recipe that will determine with a strong accuracy the popularity of an article. Our models can give an idea of the success but as Mashable has published articles making major buzz (more than 20 000 shares each; maximum reached with 840k shares), this twists our models and makes them more optimistic that they should be.

Furthermore, we could have transformed the regression problem in a classification problem, for example $nb.shares > 12000$, that could have allow to create more robust model to predict whether an article would turn viral (meaning number of shares superior to 12000) or not. This model could potentially be combined afterwards with two different models dedicated to predicting number of shares for viral and for non-viral articles exclusively.

To finish, we can keep in mind that the classification model is a lead to empower authors in defining what factors could be more impactfull to result in more engagement and virality.